

# Winter 2026 DSC 240: Introduction to Machine Learning

## Machine Problem 3

Due: Monday, Mar 9th, 11:59 pm PST

---

### Notes:

1. **This assignment is to be done individually.** You may discuss the problems at a general level with others in the class (e.g., about the concepts underlying the question, or what lecture or reading material may be relevant), but the work you turn in must be solely your own.
2. Be aware of the late policy in the course syllabus – i.e., *no late days for all the assignments*, so **it is your responsibility to turn in your assignment to Gradescope by the due time**.
3. Any updates or corrections will be posted on Piazza, so check there occasionally.
4. You can refer to online resources and cite exact references. **No copying code.**
5. You are allowed and it is recommended to use the library `sklearn` extensively. You can also use `NumPy`, `Pandas`, and any default Python libraries. **No deep learning frameworks** like `PyTorch`, `Keras`, or `TensorFlow`.
6. There is no GPU in GradeScope, so make sure your code can run on a **CPU machine with 6GB RAM**. The running time limitation in GradeScope is **30 minutes**.
7. Be sure to re-read the “**Academic Integrity**” on the course syllabus. You must complete the section below. If you answered Yes to either of the following two questions, give corresponding full details.
  - (a) *Did you receive any help whatsoever from anyone in solving this assignment?*
  - (b) *Did you give any help whatsoever to anyone in solving this assignment?*

In this project, you should implement **any classifier** for **binary classification** (not limited to linear classifiers). You are free to implement either any of the classifiers (you can directly use `sklearn` or implement the classifiers yourself) studied in class (e.g. linear logistic regression, SVM, Naive Bayes) or more complex classifiers like a multi-layer perceptron (MLP) classifier, Kernel SVM, Boosted Decision Trees, etc) The goal is to predict the decision of credit card applications. Remember to keep the function name and argument of `run_train_test(training_data, testing_data)` unchanged.

Again you can make use of **any** of the classifiers provided in `sklearn` as long as it fits into the memory and run-time constraints. A short example of how to implement any classifier from `sklearn`:

```
# import classifier
from sklearn.dummy import DummyClassifier
# create classifier object
clf = DummyClassifier(strategy="most_frequent")
# fit input data for training
clf.fit(X_train, y_train)
# make prediction
y_pred = clf.predict(X_test)
```

## Problem Overview

In this project, we focus on a real-world classification problem: whether the bank should approve the application of the credit card based on client information.

*Credit score cards are a common risk control method in the financial industry. It uses personal information and data submitted by credit card applicants to predict the probability of future defaults and credit card borrowings. The bank can decide whether to issue a credit card to the applicant. Credit scores can objectively quantify the magnitude of risk. Credit score cards are based on historical data. Given the client's historical behaviors, the bank can identify whether this client is a risky customer that tends to pay late billing. Combining this information with the corresponding application, the bank can make a better decision on credit card applications to reduce the number of risky customers. Usually, users at risk should be 3%, so the imbalanced data is a big problem.*

Change the Python 3 program `mp3.py` that creates a classifier (for example: MLP classifier or SVM or any as per your choice) for the binary classification problem. The training and development data sets are available in the starter package. The detailed data information can be found in `data.md`. The content of the starter package is:

- `data/*`: includes `train.csv` for training and `dev.csv` for development. We will have an extra private test set for the final grading.
- `data.md`: describes the data information.
- `mp3.py`: the Python file you need to work on and submit.
- `MP3.pdf`: this instruction file

## Code Instructions

In this project, you need to implement any (supported by sklearn) classifier for binary classification. As an example, you can refer to the multi-layer perceptron classifier. Specifically, there are several things to do:

1. Preprocess the features. There are different kinds of features: binary features, continuous features, and categorical features. You may conduct any form of feature selection, preprocessing, or feature expansion. Examples of data-preprocessing with sklearn can be found here: <https://scikit-learn.org/stable/modules/preprocessing.html>
2. Implement any classifier with sklearn.
3. Update the parameters of the classifier based on the training data.
4. Evaluate the performance of the development set.
5. Try different hyperparameters / preprocessing methods to get a better performance.

`run_train_test(training_data, testing_data)` function is called to train a classifier based on the training data and return the prediction on the testing data.

### Some Hints:

- You **may not need** to use all the features. There could be some redundant features.
- You **must** make sure that the training and test data are processed in the same way.
- For continuous features, you can either normalize them, or divide them into several bins and convert them to categorical features.
- Examples of some typical values of hyperparameters that are present in various classifiers (*please refer to sklearn documentation for any classifiers that you use to get exact information on what hyperparameters to tune*):
  - Hidden size: 8, 16, 32, 64, 128, ...
  - Hidden layer: 1, 2, 3, 4, ...
  - Maximum epoch: 100, 200, 300, 400, ...
  - Learning rate: 0.01, 0.001, 0.0001, 0.03, 0.003, ...
  - Batch size: 8, 16, 32, 64, 128, ...
- For the imbalanced dataset, you can refer to <https://www.kdnuggets.com/2017/06/7-techniques-handle-imbalanced-data.html>

## Evaluation Instructions

In this project, we will use the **F1 score** for evaluation. The definition can be found in previous MP2 instructions or here. Specifically, we calculate the F1 score based on the positive class, which is the risky customers with target=1. We have already implemented the evaluation metric in mp3.py.

You can test your program with the `train.csv` and `dev.csv` provided in this starter package. You can directly run `mp3.py` for checking:

```
python mp3.py
```

This will output several metrics you get on the dev dataset. The final score will be graded on another holdout private test set – different from your dev set (we will train using the same training set).

## Submission Instructions

You should upload `mp3.py` to Gradescope for grading. Please put all the code into one file.

**Note:** *Make sure that your data pre-processing code should either be implemented or called by `run_train_test(training_data, testing_data)`*

## Grading Instructions

The score on Gradescope is computed based on the **F1 score** using the following rule:

**Manual Grader (50 points):** If your code can run successfully and it trains a machine learning classifier, you will get full points. Otherwise, your score will be based on completeness. Simply "return 0" does not get you anything.

**AutoGrader (50 points):** Your score depends on the **F1 score** on the **private test dataset**:

- **15% weightage:** If your F1 score is equal to 0.40 or more, you get 15%.
- **35% weightage:** Additional prorated marks when your F1 score is above 0.40:

$$\text{Score} = 35\% \times \frac{(\text{Your score} - 0.40)}{(0.50 - 0.40)}$$

- If your F1 Score is below 0.40, Score = 0.

**Leaderboard (10 bonus points):** This score is based on your ranking on the leaderboard:

- **Top 10: +10 bonus points.** Top 10 submissions get 10 bonus points in this part.

Submissions with scores identical to the boundary will also be included. For example, if the F1 score of the 10th person is 0.41, then all submissions with F1 Score = 0.41 will be included in the Top 10.

**TA baseline F1 Score: 0.50**

**Additional Leaderboard Instructions:**

- Your leaderboard submissions will be evaluated on dev data (`dev.csv`) shared with you until the deadline. After the MP3 assignment deadline, we will re-evaluate your final submissions on private test data and use that score for the **final** leaderboard ranking.
- **Note:** your leaderboard scores may change after MP3 is completed. We will only report your dev-set score on the leaderboard before the deadline.

So if you train on the dev set you may overfit and get a very high F1 score on the leaderboard before MP3 completes, but you will see a drastic decrease in your final leaderboard F1 score.